

Does Anybody Really Know What Time It Is? Automating the Extraction of Date Scalars

The Doctor

Martha Jones

Captain Jack

c/o Time Lords
The Citadel
Gallifrey

{doctor,martha,jack}@tardis.org

ABSTRACT

With the advent of modern data visualization tools, data preparation has become a bottleneck for analytic workflows. Users who are already engaged in data analysis prefer to stay in the visualization environment for cleaning and preparation tasks whenever possible, both to preserve analytic flow and to take advantage of the visualization environment to spot inconsistent data. This has led many visualization environments to include simple data preparation functions such as scalar parsing, pattern matching and categorical binning in their analytic toolkits.

One of the most common scalar parsing tasks is extracting date and time data from string representations. Several RDBMSes include date parsing “mini-languages” to cover the wide range of possible formats. Analysis of the usage of such functions in one visualization system shows, however, that the date parsing language syntax is difficult for users to master.

In this paper, we present two algorithms for automatically deriving date format strings from a column of data, one based on minimum entropy and one based on natural language modeling. Both have similar accuracies of over 90% on a large corpus of date columns extracted from an online data repository and are in substantial agreement with each other. The minimal entropy approach is also fast enough to produce results within the user’s perceptual threshold, making it suitable for interactive work.

Categories and Subject Descriptors

D.3.3 [Data Processing]: Data Cleaning, Natural Language Processing

General Terms

Algorithms, Performance

1. INTRODUCTION

In recent years, there has been a growth of interest in data visualization technologies for human-assisted data analysis using systems such as Polaris [21] and Spotfire [5]. While computers can provide high-speed and high-volume data processing, humans have domain knowledge and the ability to process data in parallel, notably by using our visual systems. Most importantly, humans provide the definition of what is valuable in an analysis. Accordingly, human/computer analytic systems are essential to extracting knowledge of value to humans and their societies from the large amounts of data being generated today.

1.1 Interactivity

Visualization systems are most effective when they are interactive, thereby allowing a user to explore data and connect it to their domain knowledge and sense of what is important without breaking cognitive flow. In recent years, a number of such systems have been developed, both by the academic community and by the commercial sector. Exploration of data consists not only in creating visual displays, but also in creating and modifying domain-specific computations. Consequently, most data visualization systems include facilities for defining such calculations as part of the data model being analyzed. The most effective systems allow users to define these calculations as part of the analytic interaction, which permits the user to stay in the flow of analysis [18].

1.1.1 Clean as You Go

During the analytic process, a user may discover that parts of the data are not yet suitable for analysis. Solutions to this problem are often provided by data preparation tools external to the visual analysis environment, which requires the user to break their analytic flow, launch another tool and reprocess their data before returning to their analysis. If the user does not own this process (*e.g.* it is the responsibility of another department), then there can be significant delays (including “never.”) More subtly, the result of updated external processing may not be compatible with the user’s existing work, which can lead to more time lost reconciling the new data model with the existing analysis.

From the user’s perspective, the boundary between preparation and analysis is not nearly so clean cut. Bad data is often discovered using visual analysis techniques (*e.g.* histograms or scatter plots) and it is most natural for the user

to “clean what she sees” instead of switching to a second tool. This leads to an “adaptive” process whereby users will prefer leveraging existing tools in the analytics environment (no matter how well suited to the task) over switching to another application. Thus a well-designed interactive visual analysis environment will provide tools that enable users to perform such cleaning tasks as interactively as possible.

1.1.2 Function Libraries

While the existing standards for query languages are helpful for defining a core set of row-level functions, there are many useful functions beyond the standards that are only explicitly supported by a subset of RDBMSes, often with different names. And even when a database does not provide an implementation of a particular function, it is usually possible to generate it as an inline combination of existing functionality. The calculation languages of visualization systems can thus be designed to provide a *lingua franca* that tames this Babel of dialects by providing a common syntax for as many of these functions as possible.

1.2 DATEPARSE

The focus of this work is the usability of one such function from our visualization system that we will call **DATEPARSE**.

1.2.1 Scalar Dates

The SQL-99 standard defines three temporal scalar types: **DATE**, **TIMESTAMP** and **TIME**. They are typically implemented as fixed-point types containing an offset from some epoch (*e.g.* Julian Days.) This makes them compact to store using column store compression techniques such as those in C-Store [22] and MonetDB/X100 [24], and further allows some temporal operations to be implemented very efficiently using simple arithmetic. Thus from the RDBMS perspective, representing dates in scalar form provides benefits for users, both in terms of available operations and query performance.




Figure 1: Categorical Date Scalars.

Our visualization system models the first two of these SQL-99 types natively; the third (pure time) is folded into **TIMESTAMP** by appending it to a fixed date of 1899-12-30 (a convention derived from Microsoft Excel). From the analytic perspective, date types are dimensional (*i.e.* independent variables) and can be used as either *categorical* (simply ordered) or *quantitative* (ordered with a distance metric) fields.

Categorical dates have a natural hierarchy associated with them generated by calendar *binning*. The visualization in

Figure 1 shows an example of a bar chart employing binned categorical dates in a year/quarter hierarchy.




Figure 2: Quantitative Date Scalars.

Quantitative dates are typically used for time series on an axis that maps the underlying distance measure to display pixel distance. The quantitative analog to categorical binning is *truncation* whereby all the dates in a bin are represented by the first date in the bin. Truncation accurately preserves the distance semantics while enabling roll-up to coarser levels of detail. The visualization in Figure 2 shows the same sales data in a time series rolled up to the quarter level.

These types of visualizations are difficult to specify with simple categorical string representations of dates. Thus the conversion of unparsed strings to date scalars is also an important component of data modeling for analytics.

1.2.2 Parsing

This parsing of date representations is one of the oldest data preparation problems we have observed. A common version of this problem is how to convert columns of integers of the form *yyyyMMdd* to date scalars. Naïve users often solve this problem by converting the integer to a string and performing some locale-dependent string operations before casting the string back to a date. Unfortunately, this approach has a number of problems:

- String operations are notoriously slow compared to scalar operations (typically 10-100x slower in modern RDBMSes)
- The parsing code was hard to understand and maintain because it used a verbose, general-purpose string-handling syntax instead of a specialized domain language.
- Default parsing of date formats is locale-dependent, and may not work when the analysis is shared across an international organization (*e.g.* between the US and European offices)
- The parsing code is hard to understand and maintain because it uses a verbose, general-purpose string-handling syntax instead of a specialized domain language.

And this is but a single date format. Our studies of online data collections suggest that there are hundreds of distinct

temporal date formats in user data sets. Some are common, but others can be quite idiosyncratic. Table 1 shows a selection of unusual date formats found in our test data. The first example shows a time zone in the middle of the date and a year after the time; the second shows a leading unmatched bracket and a colon between the date and time components; the third shows confusion between the seconds’ decimal point and the time part delimiter; the fourth shows a two digit year apostrophe on a four digit year and the fifth shows a dash separating the date and time components.

ICU Format	Example
EEE MMM dd HH:mm:ss zzz yyyy	Fri Apr 01 02:09:27 EDT 2011
[dd/MMM/yyyy:HH:mm:ss	[10/Aug/2014:09:30:40
dd-MMM-yy hh.mm.ss.SSSSSS a	01-OCT-13 01.09.00.000000 PM
MM 'yyyy	01 '2013
MM/dd/yyyy - HH:mm	04/09/2014 - 23:47

Table 1: Unusual Date Formats.

Our first attempt to help users with this problem was to add a new function to the calculation language called `DATEPARSE`. This function would take a string column and convert it to a datetime using a special purpose domain language for describing dates. Such functions exist in a number of databases (*e.g.* MySQL, Oracle and Postgres), so it was a natural addition to the function library.

We chose the date parsing syntax defined by the International Components for Unicode (or ICU) project [4] for the common domain language because it mirrored what was already available in our code base, and translating this syntax into the formats used by various database vendors was relatively painless. There are a few patches of non-overlapping functionality, but they tend to be obscure and we provide warnings when functionality is missing.

1.2.3 Usability

After providing this new function to our user base, we waited a few months and then investigated how it was being applied in a free, cloud-based version of the system. We found that although there were a number of new calculations using `DATEPARSE`, the error rate for the domain language syntax was about 15%. `DATEPARSE` was capable of solving the problem, and some users were able to discover it, but the syntax did not appear to be easy to use reliably.

1.3 Automated Format Extraction

One solution to this problem might have been to design a graphical environment that enabled users to construct valid patterns. This would have involved a substantial development effort with no guarantee that the result would be correct if the user misunderstood the environment. Instead, we have developed two algorithms for automating the derivation of the format string, each of which uses a different machine learning technique. Both algorithms result in over 90% parsing accuracy – and 100% syntactic correctness because they are machine generated.

2. RELATED WORK

Data preparation has been considered an analytic bottleneck since at least the description of the Potter’s Wheel sys-

tem [19]. Since then, several other interactive data preparation systems have been proposed, including Data Wrangler [13] and Google Refine [12]. While effective, these systems all make assumptions about possible date formats, which we suggest are too restrictive for real world data.

Various approaches have been described for deriving regular expressions, and a good overview is provided by Li *et al.* in their paper on the ReLIE system for deriving regular expressions given a starting expression provided by a domain expert [16]. The Minimum Descriptive Length technique first described in Rissanen [20] was used in [19] to generate regular expressions.

There are several bodies of research on developing semantic parsers and grammars for interpreting time and date expressions. Lee *et al.* use a combinatory categorical grammar by combining the TIMEX3 standard [3] with contextual cues such as document creation time to determine the reference for parsing time expressions such as ‘2nd Friday of July’ [15]. Related work by Gabor *et al.* employs a probabilistic approach for learning to interpret temporal phrases [6]. Han and Lavie developed a formalism called the Time Calculus for Natural Language (TCNL), designed to capture the meaning of temporal expressions in natural language. In this formalism, each temporal expression is converted to a formula in TCNL, which then can be processed to calculate the value of a temporal expression. Their temporal focus tracking mechanism allows correct interpretation of cases like ‘I am free next week. How about Friday?’, where the TCNL formula for Friday, reflects the occurrence of ‘next’ in the preceding sentence [10]. In all this previous research using natural language techniques, the interpretation of temporal expressions are done individually, using the presence of lexical tokens such as ‘next’ and ‘past’ along with the tense of the verb tokens.

2.1 Contributions

Our work differs from these previous methods for parsing date-time expressions as we are using an entire column of data as context without the necessary presence of rich lexical identifiers to determine temporal context. This sort of data tends to be prevalent in datasets that are used in visualizations.

In particular, our contributions are:

- By analysing an on line corpus, we provide evidence that practical date parsing requires the ability to recognise hundreds of formats;
- We show how to extend prior work on Minimum Descriptive Length structure extraction to generate a freely available date format domain language with over 90% accuracy;
- We describe a second Natural Language Processing technique for generating the same date format domain language with similar accuracy. We describe how the basic algorithm is extended to support grammar variants and constraints unique to date formats. The parsing algorithm is extended to compute an overall dominant pattern over a data column;

- We describe some limitations of this domain language that would improve its utility.

2.2 Overview

The rest of this paper is organized as follows: The next section introduces the parameters of the problem space. The following two sections describe the two different algorithms, one using Minimum Descriptive Length and the other using Natural Language Processing. In section 6, we evaluate the algorithms on a corpus of 30K columns, both by sampling the outputs manually and then by using the algorithms to validate each other. We then discuss future work in section 7 and conclude in section 8.

3. PARAMETERS

Before delving into detailed descriptions of the two algorithms, we describe the common elements of the problem they are intended to solve. These are the input data, the output language and the interpretation of the output language for partial dates.

3.1 Input Data

The training data and test data are taken from data sets published to a free, online data analysis website. The published data includes both the raw data and how it was used for analysis. The service limits data sets to a maximum of 100K rows stored in a database which supports an implementation of `DATEPARSE`. We selected a large number of string columns whose names suggested that they might contain temporal data. We included words from multiple languages besides English and extracted the column's locale (actually the collation) along with the data.

3.1.1 NULL Filtering

Each column was then reduced to a maximum sample set of 32. We excluded domain values that appeared to be representations of `NULL`:

- Values containing the substring `NULL`
- Values containing no digits and at most one non-whitespace character (*e.g.* " / / ")
- Values on a list of empirically determined common `NULL` values (*e.g.* 0000-00-00, NaN)

Columns that had no remaining valid samples were discarded.

3.1.2 Sampling

The remaining samples were hashed and sorted on the hash value, with the top 32 values being retained as the column's sample. We can increase the sample size if needed, but for dates, it appears to be an adequate number. In any case, the median number of non-null rows per domain in our test data is 50, so increasing the sample size would have little benefit.

3.1.3 Numeric Timestamps

After `NULL` filtering, there remained one common class of date representation that was unsuited for parsing via our date format syntax, namely numeric timestamps. This included Unix epoch timestamps (expressed as second, millisecond or microsecond counts from 1970-01-01) and Microsoft Excel timestamps (expressed fractional days since 1900-01-01). A simple test to check that the column was numeric and inside a specific range of values representing recent dates allowed us to tag these columns to avoid analyzing them further (and incidentally to identify them for generating a simple date extraction calculation for the user.)

3.2 The ICU Date Format Language

The date format syntax we selected is the one provided by the ICU open-source project. We chose it because we were already using ICU in the code base, we had access to the source code, and it provides localized date part data for a large number of languages.

The syntax is documented at the ICU web site [4]. While fairly complete, it has a few limitations that we ran into when evaluating the algorithms:

- No support for 4 letter year abbreviations (*e.g.* Sept.)
- No support for ordinal days (*e.g.* July 4th)
- No support for quarter postfix notation (*e.g.* 2Q)
- No support for variant meridian markers (*e.g.* a.m.)

These limitations did not affect the results significantly, and in the future we hope to submit ICU extensions to handle some of these issues.

One other quirk of the ICU syntax may be a contributing factor to the user confusion around writing correct ICU date formats. The use of lexicographical case in the meta-symbols of the format language can be confusing (*e.g.* `y` is used for years, but `M` is used for months while `m` is used for minutes.) Another advantage of an automated algorithm is that it hides such problems from most users, significantly improving the usability of the function.

3.3 Partial Dates

Many of the date formats that we encountered were incomplete dates, which necessitated creating rules for what date scalar they represented.

ICU's date parsing APIs allow the specification of default values for parts when a format does not contain them. In our implementation, all time fields are set to 0 (midnight) and the date fields are set based on whether the format contains any date part specifications. When date parts are present, we use 2000-01-01 as the set of default date parts as it is the start of a leap year. When dealing with pure time formats, we model the output as a date/time and use 1899-12-30 for the date parts.

ICU will also parse Time Zones (which the algorithms recognize but the visualization system does not model) and

Quarters (which we interpret as the first month of the period.) RDBMSes such as Oracle and Postgres that support time zones will be able to take advantage of identified time zone fields.

4. MINIMAL DESCRIPTIVE LENGTH

The first algorithm is a Minimum Descriptive Length [20] approach derived from the domain system presented in the Potter’s Wheel system of Raman *et al.* [19]. We describe a number of extensions to their structure extraction system to support more complex redundancy, non-English locales, improved performance and date-specific pruning.

4.1 Domains

Potter’s Wheel presents an algorithm for deriving a common structure for a set of strings by breaking each string down into a sequence of domains. These domains are described by an interface that includes:

- A required inclusion function to test for membership in the domain (`match`)
- An optional function to compute the number of values in the domain with a given length (`cardinality`)
- An optional function to update statistics for the domain based on a given value (`updateStatistics`)
- An optional function to prevent consideration of a domain that is redundant (`isRedundantAfter`).

In our approach, we implement all of these functions, but with significant changes to the last one, which we will describe below.

With this interface, we can now define a set of domains for each date part that we wish to be able to parse. These are mostly straightforward enumerations and numeric ranges, each tagged with the ICU format code. Since the ICU parser is flexible about parsing single or double digit formats, we use double-digit formats, but accept one or two digits. One important exception to this rule is for years, which are fixed width fields (2 or 4).

We also included some enumerated domains for handling constant strings and some simple regular expression domains for delimiters such as spaces, punctuation or alphabetic characters. We found that the inclusion of arbitrary numeric domains caused the run time to grow exponentially as the number of possible matches could not be pruned intelligently. This restriction extends to domains that can contain arbitrary digit sequences (such as any). Because of this restriction, the algorithm cannot extract non-date numeric fields.

4.2 Redundancy Extensions

A difficulty in using this kind of structure extraction is that the algorithm for enumerating structures is exponential in the number of domains. This is especially true in the date format problem because there are identical domains (*e.g.* months and meridian hours), nearly identical domains (*e.g.*

days and hours) and there are often no field delimiters (*e.g.* 2012Mar06134427). To handle this, we have extended the existing pruning API with two other sets of domain identifiers:

- A set of *prunable* identifiers, which are not allowed to precede the domain. For example, once we have a month field, no other month fields should be generated. Each month domain therefore lists all the month domains in its prunable set.
- A set of *context* identifiers, one of which must have been previously generated before the domain is considered. For example, a meridian domain can only be generated once an hour field has been found, but there may be other intervening fields.

4.3 Performance

Structure enumeration is computationally expensive, so we added a number of enhancements to the original domain extraction algorithm to keep the run time low enough for interactivity.

4.3.1 Domain Characteristics

Date domains typically have small widths, so we found it advantageous to provide the shortest and longest match sizes for use in structure enumeration and matching.

Date domains are also often uniform in that adding more characters to a mismatch will not help. For example, a 2-digit day domain that does not match a 1-letter substring will not be able to generate a match by adding more characters.

4.3.2 Parallel Evaluation

We have also identified two opportunities for parallel computation during structure enumeration.

The first computationally expensive operation is the enumeration of structures for a given sample. To parallelize this step, each thread is given a subset of the samples and independently produces a set of candidate domains. When all threads have completed, the duplicates are removed to produce a single list of candidate structures.

The second computationally expensive operation is the evaluation of each generated structure over the entire sample set. This includes computation of the MDL, recording of domain statistics and parameterizing the structure. These tasks are simple to parallelize, because there are no overlaps between the data for each structure.

4.4 Unparameterization

Domain parameterization is an important part of generating compact representations via MDL, but it creates problems for date recognition. For example, if a set of dates contains a constant month string (*e.g.* all values are in September) it is important to keep track of the month name domain. Consequently, when we parameterize a constant generic <Word> domain, we tag it with the date part domain that it matches (if any). We then need to apply an additional pruning step

to remove any structures that also found an equivalent domain (*e.g.* two-digit month). These rules are equivalent to the context-based redundancy rules above, but have to be applied again after parameterization of generics.

4.5 Global Pruning

The pruning rules used for the structure extraction reduce the search space dramatically, but they are also contextual and can only look backwards. The domains also contain a fair amount of ambiguity that requires the application of domain knowledge. We therefore found it necessary to add some post-generation global pruning rules:

- The set of date parts cannot contain place value gaps (*e.g.* structures that have year and day without month are removed.)
- Similarly, the set of time parts cannot contain place value gaps and must also be in place value order (times are never written in orders such as *mhs*.)
- The existence of time parts cannot make dates incomplete (*e.g.* patterns like year-day-hour are removed.)
- Two digit years require special handling. In particular, they cannot appear adjacent to a two-digit field if the structure contains punctuation. (This can come up in some small early 21st century year domains where a two-digit year can masquerade as almost any numeric field.)

These global pruning rules are simple and intuitive, but are essential to further reducing the search space.

4.6 Locale Sensitivity

Providing an acceptable international user experience requires correctly handling the locale of the column text. Accordingly, at the start of the structure extraction we use the column locale to create a set of domains containing locale-sensitive strings such as month names. We also use the locale to map these strings to upper- and lower-case in addition to the ICU mixed case strings. This enables us to accurately compute MDL statistics without having to map the candidate strings at runtime (which would be slow).

Knowing the locale of a string is not always helpful. In our data set, we found numerous cases where the locale was specified (*e.g.* Sweden) but the data was actually in English. Accordingly, we test both locales and rank the combined results.

We have built this system for the Gregorian calendar only as our visualization system does not support non-Gregorian calendars and we have little evidence of other calendars (*e.g.* Hebrew, Islamic Civil) being used for analytics. ICU supports non-Gregorian calendars and we expect the algorithm could be extended to them as well if needed.

4.7 Ranking

MDL structure analysis naturally produces a ranked list of format candidates, but we have found that a number of other properties of the formats should be preferred over simple compactness:

- Since we have a set of samples, we can apply the candidate format to the strings to see how well it performs. Formats with fewer parse errors are preferred.
- Date parts can be considered a place-value system, so we prefer “more significant” components (*e.g.* month-day-year over hour-minute-second).
- If two formats from different locales give the same results, prefer the original column locale. The sample set may have missed an example where this could be important.
- If the format has an ambiguous date order (*e.g.* all days are less than 12), then prefer the default date order of the locale. Again, the sample set may have missed a counterexample, so this is the best option.
- Once these semantic preferences have been considered, we then prefer the more compact (MDL) representation. The output of the algorithm is now an ordered list of formats and associated locales. These can then be used to drive a user interface that allows the user to choose between the possibilities or the top-ranking format can simply be used automatically.

5. NATURAL LANGUAGE PROCESSING

The main characteristic of formal grammar is the presence of production rules that are used to rewrite an expression, and a starting symbol which is the origin of all forms of a given expression. In the production rules, we find non-terminals and terminals, with parsed expressions containing only terminals. A formal grammar is considered “context free” when its production rules can be applied regardless of the context of a nonterminal. This has several advantages. First, it allows for modular syntax definitions, which simplifies grammar development and enables reuse. Second, it grants total freedom in structuring a grammar to best fit its intended use. The power of such a grammar is its expressiveness, and is not constrained by pattern rigidities that are prevalent with regular expressions and text processing [9].

5.1 Context-Free Grammar

ICU date formats are well defined both structurally and semantically, and can be defined by a context-free grammar (CFG). A CFG consists of a set of non-terminal symbols X , a set of terminal symbols β , a start non-terminal symbol $S \in X$ from which the grammar generates the strings, and a set of production rules τ , with each rule of the form $X \rightarrow \beta$ [11].

While there are several functionally equivalent notations for representing CFG, we use the *Backus-Naur Form* (BNF) for defining the grammar rules for DATEPARSE formats. In particular, we use the *Extended Backus-Naur Form* (EBNF) as the notation is more compact and readable for frequently used constructions [9]. We use a Cocke-Younger-Kasami (CYK) parser for syntactically parsing a column of date-time strings based on the grammar [7, 23, 14]. The parser is a dynamic programming algorithm that identifies the highest probable syntactic parse trees, given the grammar and an input date-time string.

We define a *EBNF* grammar for identifying date-time strings below:

$\langle \text{TimeGrammar} \rangle ::= \langle \text{Hours} \rangle \text{'.'} \langle \text{Minutes} \rangle \text{'.'} \langle \text{Seconds} \rangle$
 $\langle \text{TimeZone} \rangle \langle \text{AMPM} \rangle$ (for 12-hour formats)

$\langle \text{DateGrammar} \rangle ::= \langle \text{BigEndianDate} \rangle$
 $\quad | \langle \text{MiddleEndianDate} \rangle$
 $\quad | \langle \text{LittleEndianDate} \rangle;$

$\langle \text{DateTimeGrammar} \rangle ::= \langle \text{DateGrammar} \rangle$
 $\quad | \langle \text{TimeGrammar} \rangle;$

$\langle \text{BigEndianDate} \rangle ::= \langle \text{Year} \rangle \langle \text{Month} \rangle \langle \text{Day} \rangle;$

$\langle \text{MiddleEndianDate} \rangle ::= \langle \text{Month} \rangle \langle \text{Day} \rangle \langle \text{Year} \rangle;$

$\langle \text{LittleEndianDate} \rangle ::= \langle \text{Day} \rangle \langle \text{Month} \rangle \langle \text{Year} \rangle;$

$\langle \text{Year} \rangle ::= \langle \text{TwoYear} \rangle | \langle \text{FourYear} \rangle;$

$\langle \text{QuarterYear} \rangle := \langle \text{Quarter} \rangle \langle \text{Year} \rangle;$

$\langle \text{Day} \rangle ::= dd$ (where $dd \in [01-31]$, depending on month/year);

$\langle \text{Month} \rangle := \langle \text{MonthWord} \rangle | \langle \text{MonthNumber} \rangle;$

$\langle \text{MonthWord} \rangle := \text{'January'} | \text{'February'} | \text{'March'} | \text{'April'} |$
 $\text{'May'} | \text{'June'} | \text{'July'} | \text{'August'} | \text{'September'} | \text{'October'}$
 $| \text{'November'} | \text{'December'};$

$\langle \text{MonthNumber} \rangle := dd$ (where $dd \in [01-12]$);

$\langle \text{DayOfWeek} \rangle ::= \text{'Monday'} | \text{'Tuesday'} | \text{'Wednesday'} | \text{'Thursday'}$
 $| \text{'Friday'} | \text{'Saturday'} | \text{'Sunday'};$

$\langle \text{Hour} \rangle ::= \langle \text{TwelveHour} \rangle | \langle \text{TwentyFourHour} \rangle;$

$\langle \text{Quarter} \rangle ::= \text{'Quarter'} dd$ (where $dd \in [01-04]$);

$\langle \text{TwoYear} \rangle ::= dd;$

$\langle \text{FourYear} \rangle := dddd;$

$\langle \text{TwelveHour} \rangle ::= dd$ (where $dd \in [00-12]$);

$\langle \text{TwentyFourHour} \rangle ::= dd$ (where $dd \in [00-23]$);

$\langle \text{AMPM} \rangle := \text{'a.m.'} | \text{'p.m.'};$

$\langle d \rangle ::= \text{'0'} | \text{'1'} | \text{'2'} | \text{'3'} | \text{'4'} | \text{'5'} | \text{'6'} | \text{'7'} | \text{'8'} | \text{'9'};$

The parsed output is then converted to the ICU Date Format Language described in Section 3.2 in order to effectively cross-validate results from the MDL algorithm.

5.2 Grammar Variants and Constraints

A classical CFG performs poorly on inflected expressions resulting from misapplication of morphological inflection and syntactic rules. Large numbers of non-terminals are necessary for representation of all variations of such features, and the parser output would consist of many parse trees. We extend the date-time grammar by adding morpho-syntactic variants. These variants allows the parser to correct for syntactic errors (*e.g.* usage of whitespace characters and punctuation marks), capitalization errors and abbreviations (*e.g.* 'Mon' for **Monday**). We use several various external corpora for such corrections [2].

The date-time grammar also includes a large number of syntactically correct but semantically invalid

date-time expressions. While we have added range restrictions to symbols such as **Hour** (1–12 for 12-hour format and 1–24 for 24-hour format), **Days** (1–7), **Month** (1–12), there are special cases that need to be accounted for. For example, there is no 'November 31, 2015', 'February 29, 2013', or 'Sunday, May 5, 1965'. November only as 30 days in any year; 2013 was not a leap year; and May 5, 1965 was a Wednesday.

While custom production rules can be added to the existing grammar to exclude such expressions, this approach is not optimal as it leads to a rather large grammar that needs to account for every single semantically valid date-time sequence of terminal symbols. Rather, we modify the existing grammar with the following additional constraints to the Day terminal symbol for excluding such expressions:

- Restriction on the distribution of 30 and 31: Months usually alternate between lengths of 30 and 31 days. We can use $x \bmod 2$ to get an alternating pattern of 1 and 0, then just add our constant base number of days:

$$\text{Day} = 30 + x \pmod{2} \quad (1)$$

where $x \in [1..12]$ for each of the 12 months.

Except for **February**, Equation 1 addresses **January** and **March** through **July**. After **July**, the pattern should skip one, and the rest of the months should follow the alternating pattern inversely.

To obtain an inverse pattern of alternating 0 and 1, we add 1 to the dividend:

$$\text{Day} = 30 + (x + 1) \pmod{2} \quad (2)$$

In Equation 2 the number of days for **August** through **December** is correct ($x \in [8..12]$), but not for the remaining months. We hence introduce a bit-masking function so that the equation is equal to 1 over the desired domain and 0 otherwise. Multiplying a term by this expression will result in the term being cancelled out outside its domain. To mask the latter piece of our function, we need an expression equal to 1 where $8 \leq x \leq 12$. Floor division by 8 works well, since $x < 16$.

Now if we substitute this expression in the $x + 1$ dividend of Equation 2, we can invert the pattern using our mask:

$$\text{Day} = 30 + (x + \lfloor \frac{x}{8} \rfloor) \pmod{2} \quad (3)$$

- **LeapYear Restriction for February:** While the above restriction applies to all months barring **February**, we also apply a constraint to the number of days for February, based on whether the year is leap year or not. For this, we define a new symbol in the grammar called **LeapYear**. If an expression containing the month ‘**February**’ or any such variant (*e.g.* ‘**Feb**’, ‘**2**’) with the day ‘**29**’ and a year, would need to resolve the **Year** symbol to be a **LeapYear**, defined as:

$$\text{Year} \pmod{4} == 0 \quad (4)$$

Equations 3 and 4 are now added as constraints to the **Days** symbol in the grammar.

5.3 Probabilistic Context-Free Grammar

Pattern-recognition problems such as parsing date and time formats initiate from observations generated by some structured stochastic process. In other words, even if the initial higher-level production rule of the grammar is known (*i.e.* date, time or date-time), there could be several directions that the parser resolves to. For example, in a date string ‘5/6/2015’, the pattern could either be **M/d/yyyy** or **d/M/yyyy**.

In the context of CFGs, probabilities have been used to define a probability distribution over a set of parse trees defined by the CFG, and are a useful method for modeling such ambiguity [8, 17]. The resulting formalism called Probabilistic Context-Free Grammar (PCFG), extends the CFG by assigning probabilities to the production rules of the grammar. During the process of parsing the date-time pattern, the probabilities are used as a filtering mechanism to rank the pattern(s) that a given string resolves to in the grammar.

The production rules are of the form $X \rightarrow \beta$ where X is the left-hand side of the rule, while β is the right-hand side.

Given a CFG grammar G , we define $\tau_G(s)$ as the set of all possible parse trees for input date-time string s . For any $X \rightarrow \beta \in \tau(G)$, the probability $p(X \rightarrow \beta) \geq 1$. In addition, $\sum_{(X \rightarrow \beta) \in \tau_G} = 1$. The parser then chooses the parse tree with the maximum probability, *i.e.* given a date-time string s , as input, we determine the highest scoring parse tree $X \rightarrow \beta$ as $\max_{(X \rightarrow \beta) \in \tau(s)} p(X \rightarrow \beta)$.

In order to bootstrap the probability weights in the PCFG, we manually set initial probability weights to various rules in the grammar as follows:

- Similar to the MDL ranking properties described in Section 3.7, rules involving the non-terminal **DateGrammar** on the left-side of the rule, are assigned prob-

ability weights higher than those rules with the non-terminal being **TimeGrammar**. In practice, assigning $p(\text{DateGrammar} \rightarrow \beta_1) = 0.9$ and $p(\text{TimeGrammar} \rightarrow \beta_2) = 0.7$ leads to optimal ranking of the parse trees.

- For the non-terminal symbol **Day**, the constraint specified is a range between 01 and 28-31, depending on month/year. However, for $dd > 12$, $p(\text{Day} \rightarrow dd) = 1.0$, and 0.5 otherwise. This assignment helps disambiguate days from months.

5.4 Supervised Learning

Having defined initial probabilistic weights to the PCFG, we employ supervised learning with a known training set of date time formats to estimate the rule probabilities. The training corpus is a set of files obtained from an online data analysis corpus, described in Section 5.

The occurrence frequencies of the rules in the correct (disambiguated) parse trees can be determined from the corpus using the maximum-likelihood parameter estimates.

$$p(X \rightarrow \beta) = \frac{\text{Count}(X \rightarrow \beta)}{\text{Count}(X)} \quad (5)$$

where $\text{Count}(X \rightarrow \beta)$ is the number of times that the rule $X \rightarrow \beta$ is seen in the parse trees τ , and $\text{Count}(X)$ is the number of times the non-terminal X is seen in τ . These frequencies can then be normalized into rule probabilities. This method produces accurate probability estimates when trained on a sufficiently large corpus of disambiguated parse trees.

5.5 Context Extensions to the Grammar

The grammar parser’s success is often limited by data ambiguity, and incomplete expressions. For example an input string ‘3/7/2005’ could either be interpreted as the 7th day of March or the 3rd day of July depending on the date format used. While a single entry cannot resolve the parser’s ambiguity, a column of data provides more context for determining the dominant pattern prevalent in the dataset.

The CKY parser is run over the same set of samples from each validation file to generate a probabilistic distribution of possible parse trees corresponding to a set of date-time patterns. As a second pass of the parser, starting from the most probable parse tree, each parse tree is then applied to the entire file’s columnar data to determine the dominant pattern. For example, another entry in the same dataset could be ‘25/3/2007’ thus increasing the probability that the format is **dd/mm/yyyy** as opposed to **mm/dd/yyyy**.

In addition to the columnar data context, the locale is also used to help maximize the probabilistic likelihood of a particular date-time pattern. The locale of the data column is used to retrieve the corresponding corpora for helping with the parsing and for correcting syntactical errors. However, similar to the locale sensitivity issue described in Section 4.6, the locale and the data may not always match. The dominant pattern is simply computed over the data column based on the ranked set of parse tree results.

6. TESTING

6.1 Data

We scanned all files published to a free, online data analysis website. We collected the contents of columns with names containing any of the following strings: Date, month, created, dt (abbreviation for date), mes (month in Spanish), datum (date in German), fecha (date in Spanish), data (date in Portuguese), and the day character used in Chinese and Japanese. Roughly 95% of the data on the website is in English, but we attempted to include non-English data that was available. Fields of any data type other than date or datetime were analyzed, including strings, integers, and floats. One file was created for each scanned column, containing the unique non-null values in the column.

The resulting set of files were de-duplicated. This handles any duplicate files from repeated publishing events. Files over 1MB were manually reviewed and those that did not contain dates removed.

Most database and spreadsheet systems already detect a limited set of date formats. For instance, typing the string "12/31/1999" into Microsoft Excel, is automatically interpreted as the date 1999-12-31. The Microsoft Jet library used to read these text files detects a few date formats as well. Any column already converted by Excel or Jet was not included in this study.

6.1.1 Training Data

Training used a subset of the online data uploaded through February 2014. There were 30,968 files in the resulting set.

6.1.2 Verification Data

Verification used another set of data from the same website, collected from May 2014 to April 2015. This set was collected similarly, but also added columns named time. There may be some duplication between the training and verification data. There were 31,546 files in the resulting verification set.

6.2 Evaluation

Once the training data was analyzed, it was grouped by date format. A sample of each produced date format was manually labeled. This allowed us to quickly skip over very common formats like MM/dd/yyyy and focus our efforts on much less common formats. The samples were judged as to whether the produced format was reasonable and were tagged with correct formats if the produced format was unreasonable.

Further manual tagging focused on the files most likely to represent dates. Many of the files in the collection are not actually dates, such as fields named "updated by" (which contain "date"). A random sample of 850 columns named exactly "date", "time" or "month" (case-insensitive) were manually judged.

6.2.1 Minimum Descriptive Length

Testing of the MDL algorithm was performed on a 24-core Dell T7610 running Windows 7 with the data stored on a 250GB SSD.

Number of Records	31,546
Error Rate	27.95%
Analysis Speed (μs)	2,245.04
Validation Speed (μs)	1.65
Median Not Null	50

Table 2: MDL Parsing Statistics.

To test the MDL algorithm, we ran it over the set of samples from each validation file to generate a ranked list of formats for the file. Each format was then applied to the entire file's data set, recording both the number of errors and the elapsed times. In cases where we generated multiple formats and the main format produced errors, we applied the second format to the unparsed strings. The summary statistics from this processing are presented in Table 2.

The analysis speed is the average time needed per sample for structure extraction. At 2.5ms, this is well below most human perceptual thresholds for a set of 32 samples, so any latency in command execution would be restricted to the ability of the underlying database to provide the samples for analysis in a timely manner. (As a column store, the database can often supply such domains without a full table scan, further improving responsiveness.)

The validation speed is the average time needed to parse a value, and provides an estimate of how fast an ICU-based implementation can process string values into scalars and works out to 620K values per core per second.




Figure 3: MDL Error Rate

The error rate reflects the fact that only about 40% of the files have an associated format that parses the non-null values without error. To examine the error rate in greater detail, we turn now to Figure 3.

On the left hand side, we can see that the algorithm found 744 distinct formats that parsed 13,424 files with no error. This is a remarkable number of valid formats and underscores the need for this kind of algorithm. Raising the error rate threshold to 5% results in about 2500 formats found in 15,000 files, or nearly half of the files.

What do these formats look like? Figure 4 shows a histogram of the 25 most common formats containing a year format code at the 5% threshold, color-coded by error rate. (A sample value is provided to the right of each bar for illustrative purposes.) The formats have also been filtered to files with at least 5 samples. Most of the samples are clearly dates with a wide range of formats (the format where the time zone is between the time and the year is surprisingly common.)




Figure 4: MDL Output

Some of the dates are clearly just numbers, but our approach is to assume that when the user tells us that the column contains dates we should find the best fit. The samples include dates from a wide range of historical sources (*e.g.* Roman pottery dates) so we have elected to defer the date identification task to the user.

6.2.2 Natural Language Processing

We implemented the grammar, training and parsing in Python 3.4.1 using the Natural Language Toolkit libraries [1] on a quad core Dell T7600 running Windows 7. The complexity of the CYK parsing algorithm is $O(n^3|G|)$, where n is

the length of the input and $|G|$ is the size of the grammar [23]. The PCFG grammar uses a set of 22 phrase-level non-terminals and 30 terminals to classify the various constituents in the corpus. Since the CYK algorithm finds the maximum likelihood solution, there are no search errors (rather probability $p = 0$, and a modification in the grammar is required to improve accuracy.

Similar to the MDL algorithm described in the previous subsection, the CYK parser is run over the same set of samples from each validation file to generate a ranked list of parse trees representing the date-time formats. As a second pass of the parser, starting from the most probable parse tree, each parse tree is then applied to the entire file's columnar data to determine the dominant pattern(s). For the second pass, we can parallelize the CYK parsing because the dependencies between reapplying the ranked list of parse trees are very sparse.

The initial average parsing speed to compute the ranked set of probable parse trees is 0.93s, while the average time taken to compute the overall dominant pattern(s) for the entire column of data is 1.4s. While the natural language parsing implementation is in Python as the NLTK package is easily configurable, we could expect greater speed-up in parsing performance by employing C/C++ CYK parsing libraries.




Figure 5: Most common date formats identified by the NLP algorithm.

Out of the 31,546 files used for testing, the NLP parser identified a dominant format for 26,534(84.11%) where $p \geq 0.5$. 1634 unique formats were identified. Figure 5 shows a histogram of the most common formats identified.

6.2.3 Cross-Checking

We compared the results of the minimum description length and natural language processing algorithms. The two algorithms match for 97.9% of our validation set. The main differences in the results are due to small differences in the implementations. The minimum description length implementation recognizes formats with leading plus and minus signs, but these are almost certainly numbers, not dates. The minimum description length implementation also recognizes the Excel date formats, but this was not supported by the natural language implementation. There was one file that contained formats such as ‘Fall 2000’ and ‘Spring 2000’, and the two algorithms picked different seasons in their format. One other file contained a set of integers that are not actually dates, and one algorithm picked MMyyyy format, while the other picked HHmmss format.

7. FUTURE WORK

While we are satisfied with the results so far, we found a number of columns that we were not able to process with these techniques. Many columns contained multiple date formats, and we would like to recognize this situation and generate predicate-based calculations to increase the accuracy of the results and thereby make the experience even more seamless for the user. Other columns contain date ranges, which we would like to handle by generating multiple calculations, possibly by combining regular expressions with date parsing.

In this paper, we have only considered the parsing of strings, but dates are often formatted as integers (*e.g.* 201507016). It is significantly faster to decompose integers into date parts using arithmetic operations (*e.g.* `mod` and `div`) than by using slower, locale-sensitive string parsing functions. Moreover, the number of possible formats is low enough that we may be able to enumerate them. Timestamp preparation from numeric representations is another form of preparation which we would like to automate.

In the course of our research, we have also identified a number of date part variants (*e.g.* ordinal dates, four-letter month abbreviations, alternate meridian markers and post-fix quarter syntax) that we would like to commend to the ICU project and maybe provide implementations of.

8. CONCLUSION

In this paper, we have described two effective algorithms for extracting date format strings from a small set of samples, one using a minimum descriptive length approach and one using natural language techniques. Both algorithms are accurate enough to be used automatically without user involvement. The MDL algorithm is also fast enough to deploy in an interactive environment.

While validating the algorithms on a large corpus, we also found that the number of distinct formats in the wild is surprisingly high, and demonstrates the wisdom of including general-purpose date parsing functions in data visualization tools, data cleaning tools and RDBMSes. In particular, it is interesting to note that the most prominent open source RDBMSes (*e.g.* MySQL and Postgres) both have a built-in version of DATEPARSE, possibly reflecting that this is a common need that gets implemented when users are empowered

to extend the function library of an RDBMS.

9. ACKNOWLEDGMENTS

We would like to thank to Douglas Adams for encouraging this work; BBC Wales for providing the resources to produce it; and Doctor Who Live, whose harrowing live demonstration of episode production in front of several thousand people inspired this project.

10. REFERENCES

- [1] Natural Language Toolkit. <http://www.nltk.org/>.
- [2] Natural Language Toolkit Corpora. http://www.nltk.org/nltk_data/.
- [3] Timex3: A Formal Specification Language for Events and Temporal Expressions. http://www.timeml.org/site/publications/timeMLdocs/timeml_1.2.1.html.
- [4] International Components for Unicode. <http://site.icu-project.org/>, July 2015.
- [5] C. Ahlberg. Spotfire: An information exploration environment. *SIGMOD Rec.*, 25(4):25–29, Dec. 1996.
- [6] G. Angeli, C. D. Manning, and D. Jurafsky. Parsing time: Learning to interpret time expressions. In *Proceedings of the 2012 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, NAACL HLT ’12, pages 446–455, Stroudsburg, PA, USA, 2012. Association for Computational Linguistics.
- [7] J. Cocke. *Programming Languages and Their Compilers: Preliminary Notes*. Courant Institute of Mathematical Sciences, New York University, 1969.
- [8] M. Collins. Head-driven statistical models for natural language parsing. *Comput. Linguist.*, 29(4):589–637, Dec. 2003.
- [9] D. Grune and C. J. H. Jacobs. *Parsing Techniques: A Practical Guide*. Ellis Horwood, Upper Saddle River, NJ, USA, 1990.
- [10] B. Han and A. Lavie. A framework for resolution of time in natural language. 3(1):11–32, Mar. 2004.
- [11] J. E. Hopcroft and J. D. Ullman. *Introduction To Automata Theory, Languages, And Computation*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 1st edition, 1990.
- [12] D. Huynh and S. Mazzocchi. Google refine. <http://code.google.com/p/google-refine/>, July 2015.
- [13] S. Kandel, A. Paepcke, J. Hellerstein, and J. Heer. Wrangler: Interactive visual specification of data transformation scripts. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI ’11, pages 3363–3372, New York, NY, USA, 2011. ACM.
- [14] T. Kasami. An efficient recognition and syntax analysis algorithm for context-free languages. Technical Report AFCRL-65-758, Air Force Cambridge Research Laboratory, Bedford, MA†, 1965.
- [15] K. Lee, Y. Artzi, J. Dodge, and L. Zettlemoyer. Context-dependent semantic parsing for time expressions. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics*, ACL 2014, June 22–27, 2014, Baltimore,

MD, USA, Volume 1: Long Papers, pages 1437–1447, 2014.

- [16] Y. Li, R. Krishnamurthy, S. Raghavan, S. Vaithyanathan, and H. V. Jagadish. Regular expression learning for information extraction. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing, EMNLP '08*, pages 21–30, Stroudsburg, PA, USA, 2008. Association for Computational Linguistics.
- [17] C. D. Manning and H. Schütze. *Foundations of Statistical Natural Language Processing*. MIT Press, Cambridge, MA, USA, 1999.
- [18] K. Morton, R. Bunker, J. Mackinlay, R. Morton, and C. Stolte. Dynamic workload driven data integration in tableau. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data, SIGMOD '12*, pages 807–816, New York, NY, USA, 2012. ACM.
- [19] V. Raman and J. M. Hellerstein. Potter’s wheel: An interactive data cleaning system. In *Proceedings of the 27th International Conference on Very Large Data Bases, VLDB '01*, pages 381–390, San Francisco, CA, USA, 2001. Morgan Kaufmann Publishers Inc.
- [20] J. Rissanen. Paper: Modeling by shortest data description. *Automatica*, 14(5):465–471, Sept. 1978.
- [21] C. Stolte, D. Tang, and P. Hanrahan. Polaris: A system for query, analysis, and visualization of multidimensional databases. *Commun. ACM*, 51(11):75–84, Nov. 2008.
- [22] M. Stonebraker, D. J. Abadi, A. Batkin, X. Chen, M. Cherniack, M. Ferreira, E. Lau, A. Lin, S. Madden, E. O’Neil, P. O’Neil, A. Rasin, N. Tran, and S. Zdonik. C-store: A column-oriented dbms. In *Proceedings of the 31st International Conference on Very Large Data Bases, VLDB '05*, pages 553–564. VLDB Endowment, 2005.
- [23] D. H. Younger. Recognition and parsing of context-free languages in time n^3 . *Information and Control*, 10(2):189–208, February 1967.
- [24] M. Zukowski, S. Heman, N. Nes, and P. Boncz. Super-scalar ram-cpu cache compression. In *Proceedings of the 22Nd International Conference on Data Engineering, ICDE '06*, pages 59–, Washington, DC, USA, 2006. IEEE Computer Society.